

0x15 Programming Language Exercise Session



EPFL

Recap

Type Safety

*A type system is a logical system that assigns a **type** to every **term** (variable, expression, function)*

- Static vs. Dynamic
- Declared vs. Inferred
- Strict vs. Relaxed

What are the Pros and Cons of these features?

Memory Safety

*Memory safety ensures that only **valid** objects are accessed **in bounds***

Spatial memory safety: objects are only accessed in bounds

Temporal memory safety: only valid objects are accessed

Automatic Storage Reclamation (Pros and Cons)

- Garbage collection
- Reference counting
- Smart pointers/borrow checking

Thread Safety

Thread-safe code only manipulates shared data structures in a manner that ensures that all threads behave properly and fulfill their design specifications without unintended interaction.

Examples of concurrency bugs

- Race Condition
- Deadlock

Synchronization primitives and language constructs

- Locks/Semaphores/Atomic variables
- Java **synchronized** keyword
- Goroutines, etc.

Sandboxing and Compartmentalization

Sandboxing means running programs (often untrusted code) in an isolated, restricted environment

Implemented by restricting access to certain library/system calls

- `seccomp` (SELinux)
- `cgroup` (Docker)

Compartmentalization breaks a complex system into small components and limit the access of entities to only what is necessary.

Recall PoLP

Demo: C vs Rust

Spatial Memory Safety

```
#include <stdio.h>

int main() {
    int index;
    int array[] = {1, 2, 3, 4, 5, 6};

    printf("input array index: ");
    scanf("%d", &index);
    printf("array[%d] = %d\n", index,
array[index]);

    return 0;
}
```

```
use scanf::scanf;

fn main() {
    let mut index = 0;
    let array = [1, 2, 3, 4, 5, 6];

    print!("input array index: ");
    scanf!("{}", index).unwrap();

    println!("array[{}] = {}", index,
array[index]);
}
```

`index` is an variable controlled by the user, what if it's ≥ 6 ...?

Spatial Memory Safety

```
● chibinz@nixos ~/c/h/demo (main)> ./spatial
input array index: 5
array[5] = 6
● chibinz@nixos ~/c/h/demo (main)> ./spatial
input array index: 6
array[6] = 25403392
⊗ chibinz@nixos ~/c/h/demo (main)> ./spatial
input array index: 10000000000000000000
fish: Job 1, './spatial' terminated by signal SIGSEGV (Address boundary error)
○ chibinz@nixos ~/c/h/demo (main) [SIGSEGV]> █
```

C does not check bounds, may read
garbage from stack, crash, or worse
(undefined behavior)

Spatial Memory Safety

```
● chibinz@nixos ~/c/h/demo (main)> ./spatial
```

```
input array index: 5
```

```
array[5] = 6
```

```
● chibinz@nixos ~/c/h/demo (main)> ./spatial
```

```
input array index: 6
```

```
array[6] = 25403392
```

```
⊗ chibinz@nixos ~/c/h/demo (main)
```

```
input array index: 1000000000000
```

```
fish: Job 1, './spatial' termin
```

```
○ chibinz@nixos ~/c/h/demo (main)
```

```
● chibinz@nixos ~/c/h/demo (main)> rust/target/debug/spatial
```

```
input array index: 5
```

```
array[5] = 6
```

```
⊗ chibinz@nixos ~/c/h/demo (main)> rust/target/debug/spatial
```

```
input array index: 6
```

```
thread 'main' panicked at src/spatial.rs:10:39:
```

```
index out of bounds: the len is 6 but the index is 6
```

```
note: run with `RUST_BACKTRACE=1` environment variable to display a backtrace
```

```
○ chibinz@nixos ~/c/h/demo (main) [101]> █
```

C does not
garbage from

(unc

Rust explicitly check bounds, and terminates
program when index is out of bound

Temporal Memory Safety

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *ptr = malloc(sizeof(int));
    *ptr = 1;
    printf("*ptr = %d\n", *ptr);
    free(ptr);
    printf("ptr freed\n");
    printf("try to access ptr again,\n");
    *ptr = %d\n", *ptr);
    return 0;
}
```

```
fn main() {
    let boxed = Box::new(1);

    println!("boxed value: {}", boxed);
    drop(boxed);

    println!("dropping boxed ptr");
    println!("access boxed after drop: {}",
boxed);
}
```

Use after free is a frequent issue in C. In Rust, we can force a free by using **drop**.

Temporal Memory Safety

```
chibinz@nixos ~/c/h/demo (main) [101]> ./temporal  
*ptr = 1  
ptr freed  
try to access ptr again, *ptr = 8728
```

Again, C assumes object is valid, and
reads garbage from heap.

Temporal Memory Safety

```
chibinz@nixos ~/c/h/demo (main) [101]> ./temporal
*ptr = 1
ptr freed
try to access ptr again, *ptr = 8728
```

Again, C assumes object is valid, and
reads garbage from

```
error[E0382]: borrow of moved value: `boxed`
  --> src/temporal.rs:10:45
   |
2  |   let boxed = Box::new(1);
   |   ----- move occurs because `boxed` has type `Box<i32>`, which does not implement the `Copy` trait
...
6  |   drop(boxed);
   |   ----- value moved here
...
10 |   println!("access boxed after drop: {}", boxed);
   |                                           ^^^^^ value borrowed here after move
```

Rust catches the issue at compile time, and
complains we cannot use a freed (moved) value

Type Safety

```
union options {
    int option_a;
    float option_b;
};

int main() {
    union options o;
    printf("input value for option_a: ");
    scanf("%d", &o.option_a);
    printf("option_a %i\n", o.option_a);
    printf("option_b interpreted as %f\n",
o.option_b);
}
```

```
enum Options {
    A(i32),
    B(f32),
}

fn main() {
    let mut i = 0;
    print!("input value for Option::A: ");
    scanf!("{}", i).unwrap();

    let o = Options::A(i);
    println!("o.A {}", o.A);
    println!("o.B {}", o.B);
}
```

Unions allow for different interpretations of underlying bits, fundamentally unsafe. Rust have enum instead.

Type Safety

```
chibinz@nixos ~/c/h/demo (main)> ./union
input value for option_a: 1
option_a 1
option_b interpreted as 0.000000
chibinz@nixos ~/c/h/demo (main)> ./union
input value for option_a: 1073741825
option_a 1073741825
option_b interpreted as 2.000000
```

Bit representation of float 2.0 in IEEE754, is the same as the 32 bit integer $1 \ll 30 \mid 1$

Type Safety

```
chibinz@nixos ~/c/h/demo (main)> ./union
input value for option_a: 1
option_a 1
option_b interpreted as 0.0000000
chibinz@nixos ~/c/h/demo (main)> ./union
input value for option_a: 1073741825
option_a 1073741825
option_b interpreted as 2.0000000
```

Bit representation of float 2.0 in IEEE754, is the same as the 32 bit integer $1 \ll 30 \mid 1$

```
error[E0609]: no field `A` on type `Options`
--> src/union.rs:16:26
16 |         println!("o.A {}", o.A);
    |                                ^ unknown field

error[E0609]: no field `B` on type `Options`
--> src/union.rs:17:26
17 |         println!("o.B {}", o.B);
    |                                ^ unknown field
```

Rust simply disallows accessing enums by field

Type Safety

```
if let Options::A(inner) = o {  
    println!("Option::A: {:?}", o);  
    println!("Option::B by explicit cast: {:?}", Options::B(inner as f32));  
}
```

```
input value for Option::A: 2  
Option::A: A(2)  
Option::B by explicit cast: B(2.0)
```

Enums are pattern matched, alternatives must be explicitly constructed
(by casting in this case)

Thread Safety: C

```
int counter = 0;
void *increment_counter(void *arg) {
    for (int i = 0; i < 100000; i++) {
        counter++; printf("thread %d: %d\n", (int) arg, counter);
    }
}
int main() {
    pthread_t t[8];
    for (int tid = 0; tid < 8; tid++) {
        pthread_create(&t[tid], NULL, increment_counter, (void *) tid);
    }
    for (int tid = 0; tid < 8; tid++) { pthread_join(t[tid], NULL); }
    printf("counter = %d\n", counter);
}
```

Eight threads incrementing a counter in C

Thread Safety: Java

```
public class Race extends Thread {
    public static int counter = 0;
    public static void main(String[] args) {
        var threads = new Thread[8];
        for (int tid = 0; tid < 8; tid++) {
            threads[tid] = new Thread(new Race());
            threads[tid].start();
        }
        for (int tid = 0; tid < 8; tid++) { threads[tid].join(); }
        System.out.println("counter = " + counter);
    }
    public void run() {
        for (int i = 0; i < 1000000; i++) { counter++; /* print counter */ }
    }
}
```

Similar implementation in Java, does it sum up to 800000?

Thread Safety

Unfortunately, neither the C nor the Java version works...

Threads are reading and writing concurrently.

Reading an old value of `counter` will cause issues...

Try marking the `run` method directly with `synchronized`, does it work?

```
thread 3: 799913  
thread 3: 799914  
thread 3: 799915  
thread 3: 799916  
thread 3: 799917  
thread 3: 799918  
thread 3: 799919  
thread 3: 799920  
thread 3: 799921  
thread 3: 799922  
thread 3: 799923  
thread 3: 799924  
thread 3: 799925  
thread 3: 799926  
thread 3: 799927  
counter = 799927
```

C

```
counter: 799970  
counter: 799971  
counter: 799972  
counter: 799973  
counter: 799974  
counter: 799975  
counter: 799976  
counter: 799977  
counter: 799978  
counter: 799979  
counter: 799980  
counter: 799981  
counter: 799982  
counter: 799983  
counter = 799983
```

Java

Thread Safety: Java `synchronized`

`synchronized` utilize **per object instance locks**.

`counter` in this case is a static class variable.

So putting `synchronized` directly before `run` has no effects, we need to lock `Race.class` instead.

Now it works!

```
synchronized(Race.class) {  
    counter++;  
}
```

```
counter: 799987  
counter: 799988  
counter: 799989  
counter: 799990  
counter: 799991  
counter: 799992  
counter: 799993  
counter: 799994  
counter: 799995  
counter: 799996  
counter: 799997  
counter: 799998  
counter: 799999  
counter: 800000  
counter = 800000
```

Thread Safety: Rust

```
static mut counter: i32 = 0;

fn increment_counter() {
    for _ in 0..100000 {
        counter += 1;
        println!("counter: {}", counter);
    }
}

fn main() {
    let mut threads = Vec::new();
    for _ in 0..8 {
        threads.push(thread::spawn(increment_counter));
    }

    for t in threads { t.join().unwrap(); }
    println!("counter = {:?}", counter);
}
```

Let's try implementing the same logic in Rust... Does it compile?

Thread Safety: Rust

```
--> src/race.rs:8:9
8 |         counter += 1;
  |         ^^^^^^^ use of mutable static
= note: mutable statics can be mutated by multiple threads: aliasing violations or data races will cause undefined behavior
error[E0133]: use of mutable static is unsafe and requires unsafe function or block
--> src/race.rs:9:33
9 |         println!("counter: {}", counter);
  |         ^^^^^^^ use of mutable static
= note: mutable statics can be mutated by multiple threads: aliasing violations or data races will cause undefined behavior
```

Rust immediately recognizes that using global variables is bad, and points out it could lead to race conditions! Let's try to fix this.

Thread Safety

```
use std::sync::atomic::{AtomicI32, Ordering};

static counter: AtomicI32 = AtomicI32::new(0);

fn increment_counter() {
    for _ in 0..100000 {
        counter.fetch_add(1, Ordering::Relaxed);
        println!("counter: {:?}", counter);
    }
}
```

Using atomic variables in Rust solves the issue.

Thread Safety

```
thread 3: 799913  
thread 3: 799914  
thread 3: 799915  
thread 3: 799916  
thread 3: 799917  
thread 3: 799918  
thread 3: 799919  
thread 3: 799920  
thread 3: 799921  
thread 3: 799922  
thread 3: 799923  
thread 3: 799924  
thread 3: 799925  
thread 3: 799926  
thread 3: 799927  
counter = 799927
```

C

```
counter: 799970  
counter: 799971  
counter: 799972  
counter: 799973  
counter: 799974  
counter: 799975  
counter: 799976  
counter: 799977  
counter: 799978  
counter: 799979  
counter: 799980  
counter: 799981  
counter: 799982  
counter: 799983  
counter = 799983
```

Java, no **synchronized**

```
counter: 799988  
counter: 799987  
counter: 799988  
counter: 799989  
counter: 799990  
counter: 799991  
counter: 799992  
counter: 799993  
counter: 799994  
counter: 799995  
counter: 799996  
counter: 799997  
counter: 799998  
counter: 799999  
counter: 800000  
counter = 800000
```

Rust, correct result!